

The Onboarding and Retention Command Center

From one new hire to a governed, self-driving people operation.

Outcome: task completion and day-90 retention, handled compliantly


Round 1 track: The Onboarding and Retention AI Employee

What changed since Round 1. Your Round 1 AI Employee is still the brain and is not thrown away. Round 2 puts a real operation around it: the agent grows to five or more Operators on Auto, and you build a coded Command Center on the template with AI Policies, AI Insights, an AI Manager, a Data Manager, and a Workbench. The bar moves from "can you decompose the problem" to "can you run a governed operation."

THE SCENARIO

Farah no longer shepherds one hire; she runs onboarding and retention for cohorts starting every week across several locations and legal entities. Work authorization and contract deadlines differ by jurisdiction and carry legal consequence if missed. IT, payroll, facilities, and security all have to complete their piece before day one actually works. A first payroll run is wrong and nobody has noticed. A pulse comment contains a sensitive disclosure that must never appear on a dashboard and must reach the right person confidentially. A manager is quietly non-responsive to every nudge for an at-risk hire. The Round 2 build runs whole cohorts under compliance and confidentiality policy.

HOW THE AI EMPLOYEE WORKS (EXAMPLE)



*The Orchestrator runs the ninety-day program as a time-aware, stateful process across a cohort, and a sensitive disclosure hard-
routes to a confidential path.*

This is one example, not a required design. Add or remove Operator Agents, change how each one works and the logic behind it, make the flow more optimal, and make sure it handles every kind of exception. Teams differ on the how; judges compare on the result.

WHAT RUNS WHERE

LAYER	RUNS ON	WHAT IT IS
The agent	Auto	One Orchestrator plus five or more Operator Agents doing the actual people-ops work. Orchestration must be built on Supervity Auto.
Command Center	Template	The live operational picture: KPIs, agent status, exception queues.
AI Policies	Template	Configurable business rules that constrain what the agent may do alone.
AI Insights	Template	Patterns, anomalies, and recommendations from the operational data.
AI Manager	Template	Chat and orchestration surface that talks to the Auto agent.
Data Manager	Template	Live registry of every connected system and its health.
Workbench	Template	The human queue where exceptions are reviewed and cleared.

Start from the template: github.com/digitamizers/AutoPilot-Template . Clone it and run `docker compose up` for the backend, database, and Command Center shell. Generate your Workflow API key at auto.supervity.ai/u/api-keys and read auto.supervity.ai/docs before wiring. Full walkthrough in the Round 2 Participant Guide.

WHAT YOU BUILD

The agent on Auto

One Orchestrator plus five or more Operators doing the real work. Extend your Round 1 build.

Command Center

Live view of onboarding task completion, day-90 retention risk, at-risk hires, provisioning SLA, engagement trend, and a cohort view. Sensitive disclosures never appear here.

AI Policies engine

At least three editable rules: escalation thresholds for an at-risk hire, compliance deadlines by jurisdiction, confidentiality routing for sensitive disclosures, nudge cadence, and risk-score thresholds.

AI Insights

Attrition-risk clusters, provisioning bottlenecks by team or resource, manager patterns, and engagement anomalies. Sensitive-disclosure alerts surface only through the confidential path.

AI Manager

A chat and orchestration surface where a person asks the operation questions and watches the Orchestrator delegate to Operators and report back.

Data Manager

A live registry of every connected system, what each is used for, and whether the connection is healthy.

The Workbench

Two queues: at-risk-hire escalations to managers and People Ops, and a strictly confidential queue for sensitive disclosures that only the right role may open.

WHAT YOU CAN ADD IN ROUND 2

Round 1 covered the core path. These are the business functions this operation is still missing, and each one is naturally its own Operator. Pick the ones that make your build a fuller people-ops AI Employee; you are not required to build all of them.

Multi-jurisdiction compliance

Track work authorisation and contract deadlines per jurisdiction, where a miss carries legal consequence.

Cross-team provisioning

Coordinate IT, payroll, facilities, and security so day one actually works for every hire.

First-payroll verification

Check the first pay cycle and flag an error before the new hire discovers it.

Learning and milestone tracking

Track required training by day and chase whatever is missing before the milestone closes.

Attrition risk scoring

Score risk from leaver history, missed milestones, and pulse sentiment together.

Manager accountability

Escalate when a manager ignores every nudge for a hire who is already at risk.

WORKING WITH THE DATA

The dataset is how you exercise your integrations, not a file to read from. Either seed the systems you connect with it (load the records into your CRM, ticket tool, database, or an Airtable base), or park it in a store like SharePoint, OneDrive, Box, or Dropbox and point an integration at it, then pull the data back through that integration. Either way the data reaches your Operators through a live integration, at least three: one channel, one system of record, and a live exception to a person.

A way in

A new hire added via a form or a row.

A system of record

Onboarding tracked in Airtable; provisioning tasks in Jira or Asana.

A human loop

A Slack nudge to the manager, and a confidential route for a sensitive disclosure, routed through the Workbench.

Note: Google integrations (Drive, Sheets, Gmail, Calendar) are in beta for this event, so use alternatives such as Airtable, Supabase, SharePoint, OneDrive, Box, Dropbox, and Outlook.

THE ROUND 2 DATA PACK

WHAT IS IN IT

Eleven tables and roughly 4,900 rows of Workday and Peakon-style data: a worker register, onboarding tasks by milestone, a provisioning feed, engagement pulses, and a manager directory, plus the Round 2 additions: Locations_Entities, Compliance_Items, Payroll_Records, Learning_Milestones, Attrition_History, and Cross_Team_Dependencies.

WHAT IS NEW SINCE ROUND 1

Six new tables: Locations_Entities, Compliance_Items, Payroll_Records, Learning_Milestones, Attrition_History, and Cross_Team_Dependencies, covering jurisdictions, legal deadlines, first pay, required training, leaver history, and the teams that gate day one. The pack keeps the Round 1 schema, so your existing build still reads it, and adds volume and harder cases. Build for the general case: a judge may ask you to run a record you did not prepare.

MANDATORY IN ROUND 2

- 1. Agent on Auto.** One Orchestrator coordinating at least five distinct Operators, with parallel, branching, or stateful behavior.
- 2. Connected Command Center.** The coded frontend is wired to the Auto agent through the backend API and shows live agent activity, not the template's demo data.
- 3. Working AI Policies.** At least three active policies that measurably change agent behavior, editable without code, with every evaluation logged.
- 4. A live human loop.** At least one real exception routed to the Workbench with full context and resolved there.
- 5. Live integrations.** At least three across two categories (one channel, one system of record), visible and healthy in the Data Manager.

DO'S AND DON'TS

Do

- Build the Operators on Auto; code the Command Center, Policies, Insights, Manager, and Data Manager around them.
- Make policies genuinely dynamic: a changed threshold changes what the agent does, live.
- Generate insights from the data the agent actually processed.
- Route real exceptions to the Workbench with full context and a recommendation.

Don't

- Rebuild the Orchestrator or Operators outside Supervity Auto.
- Ship policy or insight pages that only show the template's static demo data.
- Let the agent invent a value on a missing field; pause and escalate.
- Hardcode to these rows, skip the human loop, or stage the demo.

Bonus. Going beyond the brief earns bonus points once the stated problem is solved: extra Operators and richer downstream actions, forecasting, self-learning where a human correction at the Workbench changes future behavior, deeper auditability, meaningful use of open-source components given a real role in the build, and genuine creativity on top of the problem statement.

Not apps. Not agents. AI Employees run the Autos.